

Camera Calibration

Edge Detection

Lecture 9

Outline of Alternate Proof: Plane + Perspective Model

- ▶ Consider $Z=0$ Plane in the world

- ▶ Then,


$$\begin{bmatrix} hx \\ hy \\ h \end{bmatrix} = \begin{bmatrix} \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z=0 \\ 1 \end{bmatrix} = \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} X \\ Y \\ 1 \end{bmatrix}$$

3x4 perspective transform matrix
Note: 3rd column does not matter
because of $Z=0$

- ▶ Since in homogeneous coordinates, scale factor does not matter, the 3x3 matrix is a projective transform between the world plane and the camera image plane.

Continued...

Outline of Alternate Proof




Transformation From Camera 2 to Camera 1 will be $H_1 H_2^{-1}$ (Why?)

It will be a projective transform if projective transformation operation forms a group (Prove)



Examples of Projective Transformations



Camera Calibration

- ▶ To relate 3D world points to 2D camera points, we need to know a lot of things about the camera
 - Camera Location X, Y, Z
 - Camera orientation α, β
 - Gimbal vector $(r_1, r_2, r_3)^T$
 - Focal length f
 - Size of CCD array
 - Center of projection
- ▶ Intrinsic parameters: internal to the camera.
 - ▶ Do not change when camera is moved
- ▶ Extrinsic parameters: External to the camera.
 - ▶ Change when camera is moved



Camera Calibration

- ▶ In general, the camera model looks like:

$$C_h = AW_h,$$
$$\begin{bmatrix} C_{h1} \\ C_{h2} \\ C_{h3} \\ C_{h4} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

- ▶ Calibration is the process of finding the parameters $[a_{11} \dots a_{44}]$
- ▶ If W_h and C_h are known, then we can solve for the unknown parameters



Camera Calibration

$$x = \frac{C_{h1}}{C_{h4}},$$
$$y = \frac{C_{h2}}{C_{h4}}.$$

$$C_{h1} = a_{11}X + a_{12}Y + a_{13}Z + a_{14} = C_{h4}x$$

$$C_{h2} = a_{21}X + a_{22}Y + a_{23}Z + a_{24} = C_{h4}y$$

$$C_{h4} = a_{41}X + a_{42}Y + a_{43}Z + a_{44}$$

$$a_{11}X + a_{12}Y + a_{13}Z + a_{14} - a_{41}Xx - xa_{42}Y - xa_{43}Z - xa_{44} = 0$$

$$a_{21}X + a_{22}Y + a_{23}Z + a_{24} - a_{41}Xy - ya_{42}Y - ya_{43}Z - ya_{44} = 0.$$



Camera Calibration

$$a_{11}X + a_{12}Y + a_{13}Z + a_{14} - a_{41}Xx - xa_{42}Y - xa_{43}Z - xa_{44} = 0$$

$$a_{21}X + a_{22}Y + a_{23}Z + a_{24} - a_{41}Xy - ya_{42}Y - ya_{43}Z - ya_{44} = 0.$$

- ▶ This equation has 12 unknowns
- ▶ Each correspondence yields two equations
- ▶ If 6 correspondences are known, we can solve for the unknowns



Camera Calibration

- ▶ Separating out the knowns and the unknowns

$$\begin{bmatrix} X_1 & Y_1 & Z_1 & 1 & 0 & 0 & 0 & 0 & -x_1 X_1 & -x_1 Y_1 & -x_1 Z_1 & -x_1 \\ X_2 & Y_2 & Z_2 & 1 & 0 & 0 & 0 & 0 & -x_2 X_2 & -x_2 Y_2 & -x_2 Z_2 & -x_2 \\ \vdots & & & & & & & & & & & \\ X_n & Y_n & Z_n & 1 & 0 & 0 & 0 & 0 & -x_n X_n & -x_n Y_n & -x_n Z_n & -x_n \\ 0 & 0 & 0 & 0 & X_1 & Y_1 & Z_1 & 1 & -y_1 X_1 & -y_1 Y_1 & -y_1 Z_1 & -y_1 \\ 0 & 0 & 0 & 0 & X_2 & Y_2 & Z_2 & 1 & -y_2 X_2 & -y_2 Y_2 & -y_2 Z_2 & -y_2 \\ \vdots & & & & & & & & & & & \\ 0 & 0 & 0 & 0 & X_n & Y_n & Z_n & 1 & -y_n X_n & -y_n Y_n & -y_n Z_n & -y_n \end{bmatrix} \begin{bmatrix} a_{11} \\ a_{12} \\ a_{13} \\ a_{14} \\ a_{21} \\ a_{22} \\ a_{23} \\ a_{24} \\ a_{41} \\ a_{42} \\ a_{43} \\ a_{44} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

or

$$CP = 0,$$



Camera Calibration

- ▶ This system $CP = 0$ is a homogeneous system.
- ▶ C is rank deficient: $\text{rank}(C) = 11$
- ▶ Has multiple solutions (other than the trivial solution)...
Can be solved uniquely only up to a scale factor
- ▶ Solution?



Solving for P

- ▶ The null space of C represents the P which are the solutions to the system $CP = 0$
 - ▶ How to find null space?
 1. `null(C)` in MATLAB
 2. Take SVD of C , as $\text{svd}(C) = USV^T$. The column of V corresponding to the eigen value of zero represents the solution
(in practice, you will have to take the smallest eigen value)
-



End of Module 1

- | | |
|--|--|
| ▶ Introduction, overview of the course, policies | ▶ 2D Transformations |
| ▶ General introduction of Computer Vision and the course modules | ▶ 2D Displacement Models <ul style="list-style-type: none">▶ Affine, Projective▶ Recovering best transformation |
| ▶ Human eye | ▶ 3D Transformations |
| ▶ Digital images | ▶ Perspective, orthographic transforms |
| ▶ Capturing color images | ▶ Camera Model |
| ▶ Image histogram | ▶ Planarity assumption |
| ▶ Interlacing | ▶ Camera Calibration |
-



Edge Detection

Applying Masks to Images

- ▶ Convolution Operation
- ▶ Mask
 - ▶ Set of pixel positions and weights
 - ▶ Origin of mask

1	1	1
1	1	1
1	1	1

1	2	1
2	4	2
1	2	1

1
1
1
1
1

Applying Masks to Images

- ▶ $I_1 \otimes \text{mask} = I_2$
- ▶ Convention: I_2 is the same size as I_1
- ▶ Mask Application:
 - ▶ For each pixel
 - ▶ Place mask origin on top of pixel
 - ▶ Multiply each weight with pixel under it
 - ▶ Sum the result and put in location of the pixel



Applying Masks to Images

40	40	40	80	80	80
40	40	40	80	80	80
40	$\frac{1}{9}$ 40	$\frac{1}{9}$ 40	$\frac{1}{9}$ 80	80	80
40	$\frac{1}{9}$ 40	$\frac{1}{9}$ 40	$\frac{1}{9}$ 80	80	80
40	$\frac{1}{9}$ 40	$\frac{1}{9}$ 40	$\frac{1}{9}$ 80	80	80
40	40	40	80	80	80

$\frac{1}{9} \times$

1	1	1
1	1	1
1	1	1

$$6 * (\frac{1}{9} * 40) + 3 * (\frac{1}{9} * 80) = 53$$

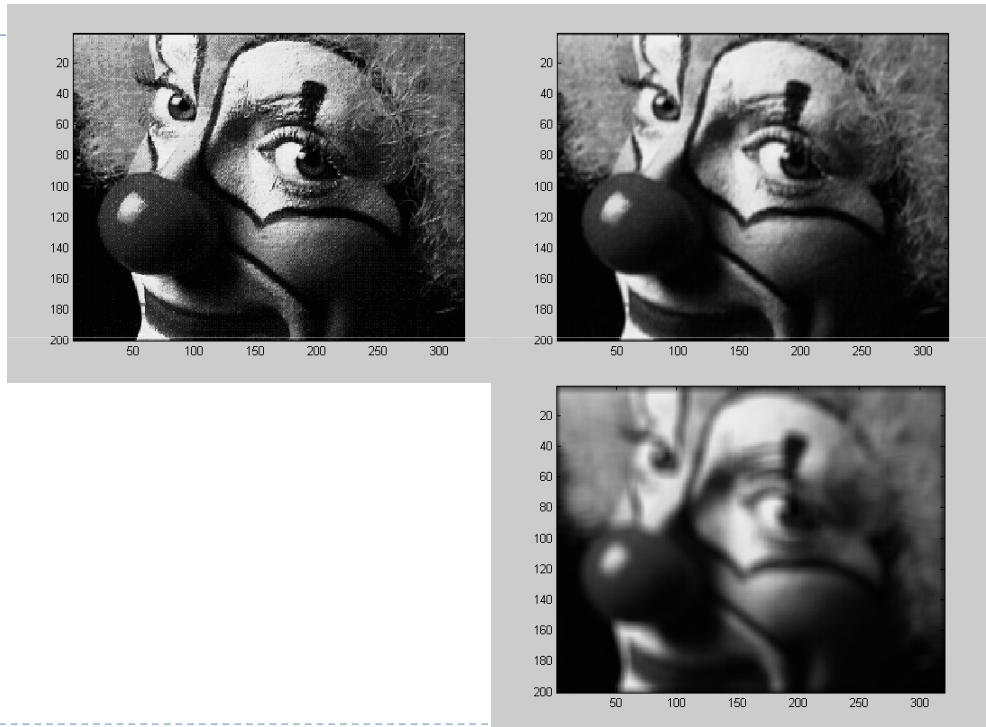



Applying Masks to Images

40	40	53	67	80	80
40	40	53	67	80	80
40	40	53	67	80	80
40	40	53	67	80	80
40	40	53	67	80	80
40	40	53	67	80	80

► Overall effect of this mask?

► Smoothing filter




What about corner pixels

- ▶ Expand image with virtual pixels
- ▶ Options
 - ▶ Fill with a particular value, e.g. zeros
 - ▶ Fill with nearest pixel value
- ▶ Or just ignore them





Edge Detection in Images

- Finding the contour of objects in a scene



Edge Detection




Choice of Mask?

- Should give 0 output on smooth regions
- Should give high output on non-smooth regions

-1	0	1
-1	0	1
-1	0	1

Prewitt mask



Why Edge Detection?

- ▶ Historical Background
 - ▶ Discarding Useless Data?
 - ▶ Simplification of Matching Problem
 - ▶ Illumination Invariant Algorithms
-

Stages in Edge Detection

- ▶ **Filtering**
 - ▶ Removal of noise
- ▶ **Differentiation**
 - ▶ Assigns high values to regions of intensity change
- ▶ **Detection**
 - ▶ Indicates regions where intensity changes are significant



Filtering

1/9 x

1	1	1
1	1	1
1	1	1

▶ Mean Filter

$$\begin{aligned}h(x, y) = & f(x-1, y-1)\frac{1}{9} + f(x-1, y)\frac{1}{9} + f(x-1, y+1)\frac{1}{9} + f(x, y-1)\frac{1}{9} \\ & + f(x, y)\frac{1}{9} + f(x, y+1)\frac{1}{9} + f(x+1, y-1)\frac{1}{9} + f(x+1, y)\frac{1}{9} + f(x+1, y+1)\frac{1}{9}\end{aligned}$$

$$\begin{aligned}h(x, y) = & f(x-1, y-1)g(-1, -1) + f(x-1, y)g(-1, 0) + f(x-1, y+1)g(-1, 1) \\ & + f(x, y-1)g(0, -1) + f(x, y)g(0, 0) + f(x, y+1)g(0, 1) \\ & + f(x+1, y-1)g(1, -1) + f(x+1, y)g(1, 0) + f(x+1, y+1)g(1, 1)\end{aligned}$$

$$h(x, y) = \sum_{i=-1}^{i=1} \sum_{j=-1}^{j=1} f(x+i, y+j)g(i, j)$$

$$h(x, y) = f(x, y) * g(x, y).$$



Filtering

- ▶ Gaussian Filter
- ▶ Pixel weight is inversely proportional to distance from origin

$$g(x, y) = e^{-\frac{(x^2 + y^2)}{2\sigma^2}}$$

- ▶ Value of σ has to be specified
-

Gaussian Filtering

- ▶ Example

$$g(x, y) = e^{-\frac{(x^2 + y^2)}{2\sigma^2}}$$

- ▶ $x = -1, y = -1, \sigma = 1$

$$g(-1, -1) = e^{-\frac{(-1)^2 + (-1)^2}{2}} \approx 0.367$$

0.367		

Gaussian Filtering

- ▶ **Implementation problem**
 - ▶ Float multiplications are slow
- ▶ **Solution**
 - ▶ Multiply mask with 255, round to nearest integer
 - ▶ Scale answer by sum of all weights



Gaussian Filtering

94	155	94
155	255	155
94	155	94

 \approx

1	2	1
2	4	2
1	2	1

Round(g/70)

5	21	35	21	5
21	94	155	94	21
35	155	255	155	35
21	94	155	94	21
5	21	35	21	5



Gaussian Filtering

0	0	0	0	1	2	2	2	1	0	0	0	0
0	0	1	3	6	9	11	9	6	3	1	0	0
0	1	4	11	20	30	34	30	20	11	4	1	0
0	3	11	26	50	73	82	73	50	26	11	3	0
1	6	20	50	93	136	154	136	93	50	20	6	1
2	9	30	73	136	198	225	198	136	73	30	9	2
2	11	34	82	154	225	255	225	154	82	34	11	2
2	9	30	73	136	198	225	198	136	73	30	9	2
1	6	20	50	93	136	154	136	93	50	20	6	1
0	3	11	26	50	73	82	73	50	26	11	3	0
0	1	4	11	20	30	34	30	20	11	4	1	0
0	0	1	3	6	9	11	9	6	3	1	0	0
0	0	0	0	1	2	2	2	1	0	0	0	0

Figure 2.3: Gaussian mask with $\sigma = 2$.




Continuous Gaussian Function

► 1-D Gaussian


$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

- Non zero from $-\infty$ to ∞
- Width is controlled by σ
- Symmetric
- Area under curve = 1

$$\int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}} = 1$$




Continuous Gaussian Function




2-D Gaussian Function


$$g(x, y) = \frac{1}{2\pi\sigma} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$




Gaussian Filter



Gaussian Vs Average



Gaussian Smoothing




Smoothing by Averaging


Is this a realistic comparison? What issues need to be looked at?




Noise Filtering



Gaussian Noise




After Averaging




After Gaussian Smoothing


Noise Filtering



Salt & Pepper Noise



After Averaging



After Gaussian Smoothing

Stage 2: Differentiation

► Continuous form

$$f' = \frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x}$$

► Discrete form (put $\Delta x = 1$)

$$f' = \frac{df}{dx} = f(x) - f(x-1)$$



Discrete Derivatives

$f(x)$	40	40	40	40	40	80	80	80	80	80
--------	----	----	----	----	----	----	----	----	----	----

$f'(x)$	0	0	0	0	0	40	0	0	0	0
---------	---	---	---	---	---	----	---	---	---	---

If $g(x)$


-1	1
----	---


 then $f'(x) = f(x) \otimes g(x)$

↑
origin


$f''(x)$	0	0	0	0	0	40	-40	0	0	0
----------	---	---	---	---	---	----	-----	---	---	---







Effect of Noise (without filtering)



Recap: Filtering Masks

- ▶ Mean Filter
- ▶ Gaussian Filter

1/9 x

1	1	1
1	1	1
1	1	1

$$g(x, y) = ce^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

5	21	35	21	5
21	94	155	94	21
35	155	255	155	35
21	94	155	94	21
5	21	35	21	5



Recap: Derivative Masks

-1	0	1
-1	0	1
-1	0	1

-1	-1	-1
0	0	0
1	1	1

-1	1
----	---

-1
1



Properties of Masks

► Filtering Masks

- All values are +ve
- Sum to 1
- Output on smooth region is unchanged
- Blurs areas of high contrast
- Larger mask -> more smoothing

► Derivative Masks

- opposite signs
- Sum to zero
- Output on smooth region is zero
- Gives high output in areas of high contrast
- Larger mask -> more edges detected



Derivative Masks

► Prewitt Operator

-1	0	1	-1	-1	-1
-1	0	1	0	0	0
-1	0	1	1	1	1



► Sobel Operator

-1	0	1	-1	-2	-1
-2	0	2	0	0	0
-1	0	1	1	2	1

► Robert's Operator

1	0	0	1
0	-1	-1	0



Application of 2-D Masks


- If f_x is derivative in x-direction, f_y is derivative in y-direction
- Gradient Magnitude

$$M = \sqrt{f_x^2 + f_y^2}$$

- Gradient Direction

$$\theta = \arctan \frac{f_y}{f_x}$$





Detection Stage - Threshold

- ▶ Gradient Magnitude

$$M(x, y) = \sqrt{f_x^2(x, y) + f_y^2(x, y)}$$


- ▶ Gradient Magnitude normalized b/w 0-100

$$N(x, y) = \frac{M(x, y)}{\max_{i=1, \dots, n, j=1, \dots, n} M(i, j)} \times 100.$$


- ▶ Application of a threshold

$$E(x, y) = \begin{cases} 1 & \text{if } N(x, y) > T \\ 0 & \text{otherwise} \end{cases}$$


Gradient Magnitude




Log of M to
enhance
visibility



$T = 10$




$T = 20$



$T = 40$





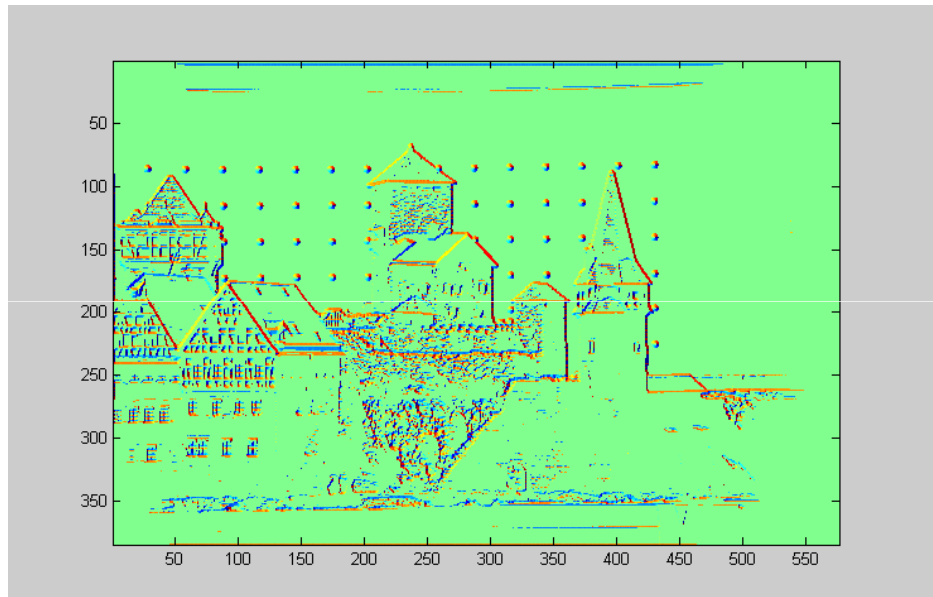
$T = 80$



What about Gradient Direction?

- ▶ Gradient Direction is always perpendicular to edge
- ▶ Direction of most change of gray levels
- ▶ Thick edges can be eliminated using gradient direction
- ▶ Weak edges also captured in this manner






Canny's Edge Detector

- ▶ **Filtering + Derivative:**
 - ▶ Uses first derivative Gaussian masks
- ▶ **Detection:**
 - ▶ Uses Non-Maxima Suppression
 - ▶ Uses Hysteresis Thresholding



First Derivative of Gaussian

- ▶ Expression?
- ▶ Effect?
- ▶ Filtering + Derivative



Canny Edge Operator




$$\Delta S = \Delta(G_\sigma * I) = \Delta G_\sigma * I$$

$$\Delta G_\sigma = \begin{bmatrix} \frac{\partial G_\sigma}{\partial x} & \frac{\partial G_\sigma}{\partial y} \end{bmatrix}^T$$


$$\Delta S = \begin{bmatrix} \frac{\partial G_\sigma}{\partial x} * I & \frac{\partial G_\sigma}{\partial y} * I \end{bmatrix}^T$$

$$f_x(x, y) = f(x, y) * \left(\frac{-x}{\sigma^2} \right) e^{\frac{-(x^2+y^2)}{2\sigma^2}},$$


$$f_y(x, y) = f(x, y) * \left(\frac{-y}{\sigma^2} \right) e^{\frac{-(x^2+y^2)}{2\sigma^2}}$$





f_x



f_y





$$M = \sqrt{f_x^2 + f_y^2}$$



$$M \geq \text{Threshold} = 10$$



Non-Maximum Suppression



Remove all points along the gradient direction that are not maximum points




Non-Maxima Suppression

$$M(x, y) = \begin{cases} M(x, y) & \text{if } M(x, y) > M(x', y') \text{ and} \\ & \text{if } M(x, y) > M(x'', y'') \\ 0 & \text{otherwise} \end{cases}$$

Where $M(x', y')$ and $M(x'', y'')$ are gradient magnitudes on both sides of edge at (x, y) in the gradient magnitude direction




Non-Maxima Suppression






Quantization of Gradient Direction

$$\theta = \arctan \frac{f_y}{f_x}$$



Non-Maximum Suppression




Hysteresis Thresholding

- ▶ Two thresholds, T_H and T_L
 - ▶ Apply non-maxima suppression to M (gradient magnitude)
 - ▶ Scan image from left to right, top to bottom
 - ▶ If $M(x,y)$ is above T_H mark it as edge
 - ▶ Recursively look at neighbors; if gradient magnitude is above T_L mark it as edge
-



Hysteresis Thresholding



Algorithm Summary

- ▶ Compute gradient of image $f(x,y)$ by convolving with first derivative of Gaussian in x and y directions

$$f_x(x, y) = f(x, y) * \left(\frac{-x}{\sigma^2} \right) e^{\frac{-(x^2+y^2)}{2\sigma^2}},$$

$$f_y(x, y) = f(x, y) * \left(\frac{-y}{\sigma^2} \right) e^{\frac{-(x^2+y^2)}{2\sigma^2}}$$



Algorithm Summary

- ▶ Compute gradient magnitude and direction at each pixel
- ▶ Perform non-maxima suppression
 - ▶ Find gradient direction at pixel
 - ▶ Quantize it in 8 directions
 $\{0, 45, 90, 135, \dots, 315\}$
 - ▶ Compare current value of M with two neighbors in appropriate direction
 - ▶ If maximum, keep it, otherwise make it zero




Algorithm Summary

- ▶ **Perform Hysteresis thresholding**
 - ▶ Scan image from left to right, top to bottom
 - ▶ If $M(x,y)$ is above T_H mark it as edge
 - ▶ Recursively look at neighbors; if gradient magnitude is above T_L mark it as edge



Choice of Sigma



Choice of Sigma

- ▶ sHalf is computed by finding the point on the curve where the Gaussian value drops below T
- ▶ $\exp(-x^2/(2*\sigma^2)) = T$
- ▶ $sHalf = \text{round}(\sqrt{-\log(T) * 2 * \sigma^2})$
- ▶ mask size is then $2*sHalf + 1$




Choice of Signma


Sigma	Size of Mask
0.5	3x3
1	5x5
2	9x9
3	13x13
4	19x19



Conv. Results



sigma = 0.5



sigma = 2

Gradient Magnitude Results



sigma = 0.5



sigma = 2




M when $\sigma = 2$



Non-maxima
suppressed image

Gradient Direction Results



Raw

Quantized



Th = 50
Tl = 10



Th = 100
Tl = 10
(incr. in Th only)



Th = 50
Tl = 40
(incr. in Tl only)

Recursion Movie

